

# Prediction of the Age of Abalone

## Introduction

In this assignment, the target is to generate an unsupervised model of to predict the age of abalone based on the physical measurement. The problem is coming from the website <https://www.kaggle.com/datasets/farkhod77/abalone-age-predict>

```
In [ ]: import matplotlib.pyplot as plt
import os
import natsort
import pandas as pd
import numpy as np
import seaborn as sns
%matplotlib inline
```

### 1. A glance of the data

To have an initial impression of the data, load the data from csv file, print the basic information. In the data frame, the number of rings represent the age. The task is to build a model to use the other physical measurements to predict the rings(age).

```
In [ ]: input = pd.read_csv("data/abalone.data.csv")
print(input)
input.describe()
```

	gender	Length	Diameter	Height	Whole weight	Shucked weight	\
0	M	0.455	0.365	0.095	0.5140	0.2245	
1	M	0.350	0.265	0.090	0.2255	0.0995	
2	F	0.530	0.420	0.135	0.6770	0.2565	
3	M	0.440	0.365	0.125	0.5160	0.2155	
4	I	0.330	0.255	0.080	0.2050	0.0895	
...	...	...	...	...	...	...	
4172	F	0.565	0.450	0.165	0.8870	0.3700	
4173	M	0.590	0.440	0.135	0.9660	0.4390	
4174	M	0.600	0.475	0.205	1.1760	0.5255	
4175	F	0.625	0.485	0.150	1.0945	0.5310	
4176	M	0.710	0.555	0.195	1.9485	0.9455	

	Viscera weight	Shell weight	Rings
0	0.1010	0.1500	15
1	0.0485	0.0700	7
2	0.1415	0.2100	9
3	0.1140	0.1550	10
4	0.0395	0.0550	7
...	...	...	...
4172	0.2390	0.2490	11
4173	0.2145	0.2605	10
4174	0.2875	0.3080	9
4175	0.2610	0.2960	10
4176	0.3765	0.4950	12

[4177 rows x 9 columns]

Out[ ]:

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell w
<b>count</b>	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.00
<b>mean</b>	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.23
<b>std</b>	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.13
<b>min</b>	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.00
<b>25%</b>	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.13
<b>50%</b>	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.23
<b>75%</b>	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.32
<b>max</b>	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.00

Considering normally the different gender will have different pattern, to simplify the calculation, split the data to male and female respectively.

```
In [ ]: df_m = input[input['gender']=='M'].drop('gender',axis=1)
df_f = input[input['gender']=='F'].drop('gender',axis=1)

print(df_m.head())
print(df_f.head())
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	\
0	0.455	0.365	0.095	0.5140	0.2245	0.1010	
1	0.350	0.265	0.090	0.2255	0.0995	0.0485	
3	0.440	0.365	0.125	0.5160	0.2155	0.1140	
8	0.475	0.370	0.125	0.5095	0.2165	0.1125	
11	0.430	0.350	0.110	0.4060	0.1675	0.0810	

	Shell weight	Rings
0	0.150	15
1	0.070	7
3	0.155	10
8	0.165	9
11	0.135	10

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	\
2	0.530	0.420	0.135	0.6770	0.2565	0.1415	
6	0.530	0.415	0.150	0.7775	0.2370	0.1415	
7	0.545	0.425	0.125	0.7680	0.2940	0.1495	
9	0.550	0.440	0.150	0.8945	0.3145	0.1510	
10	0.525	0.380	0.140	0.6065	0.1940	0.1475	

	Shell weight	Rings
2	0.21	9
6	0.33	20
7	0.26	16
9	0.32	19
10	0.21	14

## 2. Split the dataset to training/testing set

```
In [ ]: from sklearn.model_selection import train_test_split
X_train_m, X_test_m, y_train_m, y_test_m = train_test_split(
    df_m[df_m.columns.difference(['Rings'])],
    df_m['Rings'],
    train_size=0.75
)
X_train_f, X_test_f, y_train_f, y_test_f = train_test_split(
    df_f[df_f.columns.difference(['Rings'])],
    df_f['Rings'],
    train_size=0.75
)

ages_m = y_train_m.value_counts()
ages_f = y_train_f.value_counts()
```

To have better impression of the data, have a diagram to visualize the distribution of the ages and the potential pattern between ages and other biometrics.

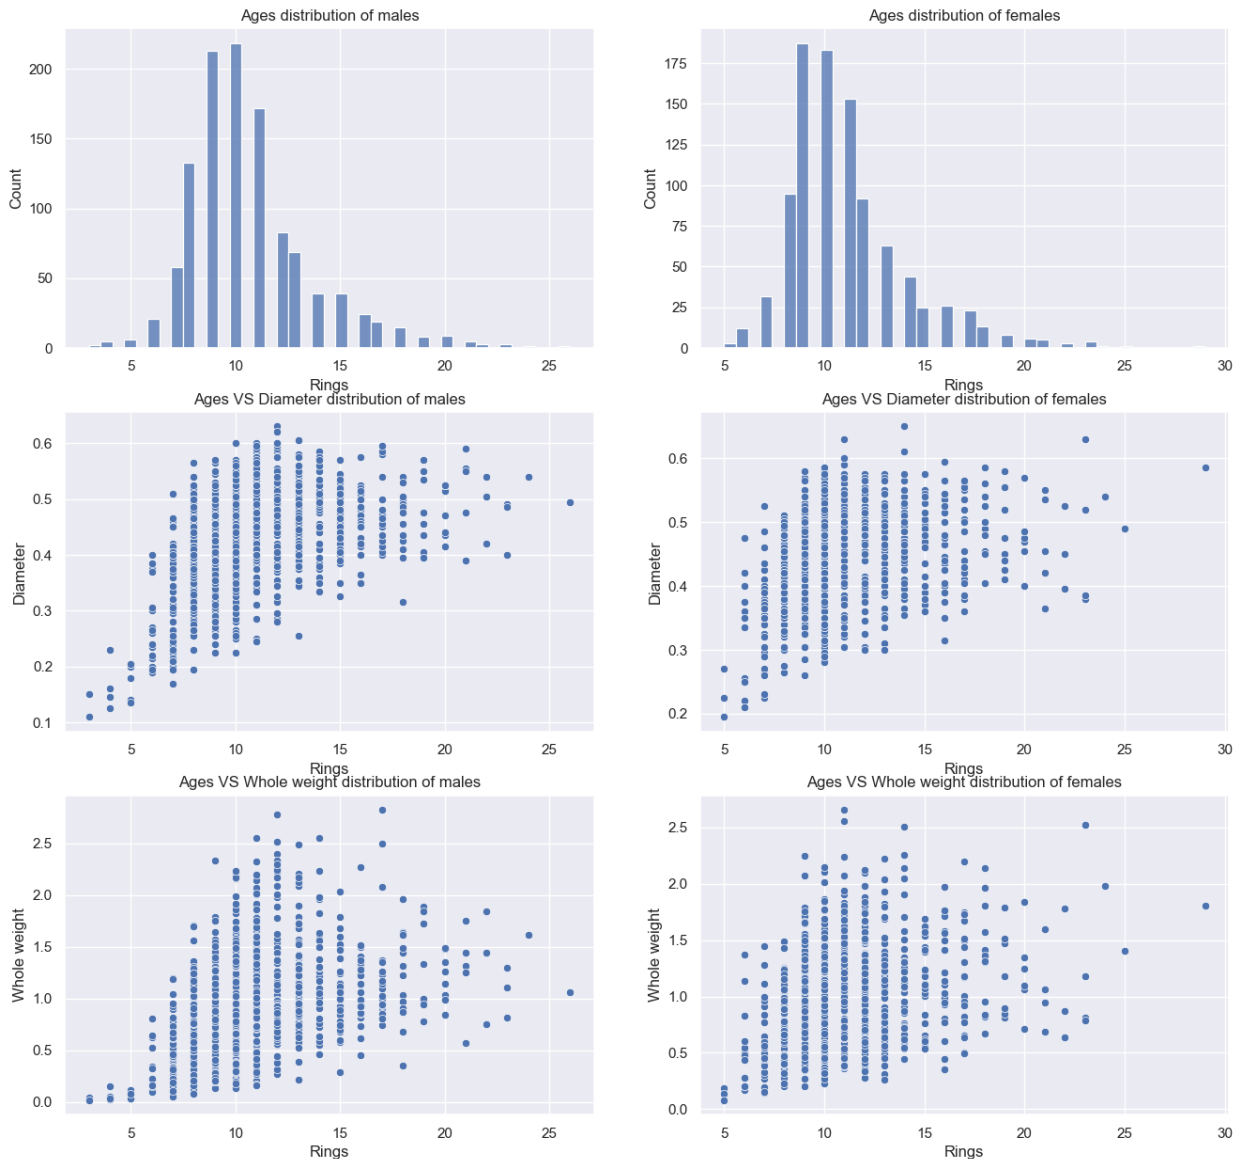
```
In [ ]: def dump_ages():
    fig, ax = plt.subplots(3,2,figsize = (16, 15))
    sns.histplot(y_train_m, ax=ax[0][0])
    sns.scatterplot(y = X_train_m['Diameter'], x = y_train_m, ax=ax[1][0])
    sns.scatterplot(y = X_train_m['Whole weight'], x = y_train_m, ax=ax[2][0])
    ax[0][0].set_title("Ages distribution of males")
    ax[1][0].set_title("Ages VS Diameter distribution of males")
    ax[2][0].set_title("Ages VS Whole weight distribution of males")
```

```

sns.histplot(y_train_f, ax=ax[0][1])
sns.scatterplot(y = X_train_f['Diameter'], x = y_train_f, ax=ax[1][1])
sns.scatterplot(y = X_train_f['Whole weight'], x = y_train_f, ax=ax[2][1])
ax[0][1].set_title("Ages distribution of females")
ax[1][1].set_title("Ages VS Diameter distribution of females")
ax[2][1].set_title("Ages VS Whole weight distribution of females")
plt.show()

```

dump\_ages()



We can get some information from the diagrams

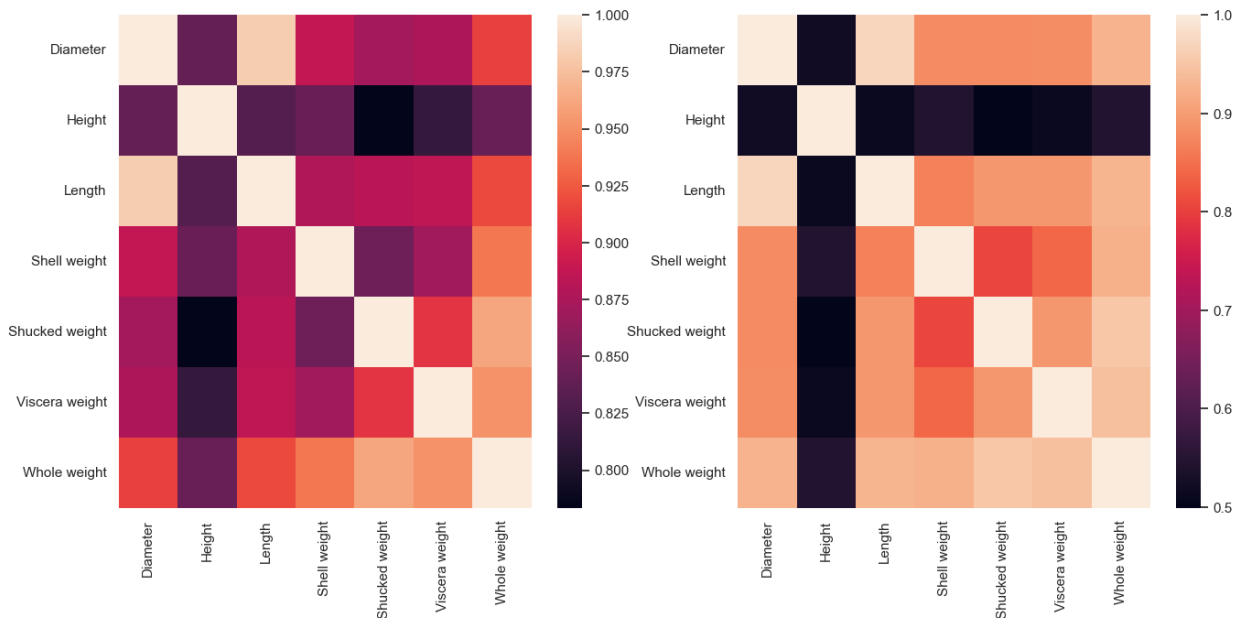
1. Most of the abalones are in the ages between 6 and 15, the number of that are younger than 6 and elder than 15 is not much.
2. The size and the weight, both are distributed widely, they are not very tight to the ages. In fact, the two factors have some kind of dependency. To have further look at the correlations, check the correlation matrix.
3. Group the data

Since the number of data may not be sufficient to be used to predict this kind of prediction (some ages may only have 1 or 2 instances), instead of predict the exact age, I'm grouping the age to fewer groups, the model will be used to predict the age group.

4 groups are defined for the purpose as below: 1->[0,8], 2->(8,10], 3->(10,12], 4->(12,]

```
In [ ]: corr_m = X_train_m.corr()
corr_f = X_train_f.corr()

fig, ax = plt.subplots(1,2,figsize = (16, 7))
sns.heatmap(corr_m, ax=ax[0])
sns.heatmap(corr_f, ax=ax[1])
plt.show()
```



From the correlation map, the association between the features is a bit high.

```
In [ ]: def age2group(df):
    top = df.max()
    df[df<=8] = top+1
    df[df<=10] = top+2
    df[df<=12] = top+3
    df[df<=top] = top+4
    df = df-top
    return df

y_train_m_group = age2group(y_train_m.copy())
y_train_f_group = age2group(y_train_f.copy())
y_test_m_group = age2group(y_test_m.copy())
y_test_f_group = age2group(y_test_f.copy())
```

Check the distribution after grouping further.


```
In [ ]: def dump_groups():
    fig, ax = plt.subplots(3,2,figsize = (16, 15))
    sns.histplot(y_train_m_group, ax=ax[0][0])
    sns.scatterplot(y = X_train_m['Diameter'], x = y_train_m_group, ax=ax[1][0])
```

```

sns.scatterplot(y = X_train_m['Whole weight'], x = y_train_m_group, ax=ax[2][0])
ax[0][0].set_title("Ages distribution of males")
ax[1][0].set_title("Ages VS Diameter distribution of males")
ax[2][0].set_title("Ages VS Whole weight distribution of males")
sns.histplot(y_train_f_group, ax=ax[0][1])
sns.scatterplot(y = X_train_f['Diameter'], x = y_train_f_group, ax=ax[1][1])
sns.scatterplot(y = X_train_f['Whole weight'], x = y_train_f_group, ax=ax[2][1])
ax[0][1].set_title("Ages distribution of females")
ax[1][1].set_title("Ages VS Diameter distribution of females")
ax[2][1].set_title("Ages VS Whole weight distribution of females")
plt.show()

```

dump\_groups()



### 3. Build Classify model

Base on the initial analysis, there are some dependent features in the original dataset, so we will apply PCA model trying to extract the primary features firstly, then a prediction model with KNN clustering algorithm is applied to build the final classifier.

```

In [ ]: from sklearn.cluster import KMeans
        from sklearn.decomposition import PCA
        from sklearn.metrics import accuracy_score, confusion_matrix
        import itertools
        class AgeGroupPredictorKNN:
            def __init__(self):
                self.n_components = 5
                self.max_iter = 10000
                self.n_init = 100

                self.pca = None
                self.knn = None
                self.X_pca = None
                self.label2group = []
            def with_params(self, params):
                if params == None:
                    return self
                for k in params:
                    match k:
                        case 'n_components':
                            self.n_components = params[k]
                        case 'n_init':
                            self.n_init = params[k]
                        case 'max_iter':
                            self.max_iter = params[k]
                return self
            def fit(self,X,y):
                self.pca = PCA(n_components=self.n_components).fit(X)
                self.X_pca = self.pca.transform(X)

                n_clusters = len(y.unique())
                self.knn = KMeans(
                    n_clusters=n_clusters,
                    n_init=self.n_init,
                    max_iter=self.max_iter,
                ).fit(self.X_pca)

                self.map_label_to_group(y)

            def map_label_to_group(self,y):
                groups = itertools.permutations(y.unique())
                best_score = 0
                l = self.knn.labels_
                for g in groups:
                    yp = [g[l[i]] for i in range(len(l))]
                    sc = self.score(y, yp)
                    if (sc > best_score):
                        best_score = sc
                        self.label2group = g
            def predict(self, X):
                x_pca = self.pca.transform(X)
                l = self.knn.predict(x_pca)
                g = self.label2group
                yp = [g[l[i]] for i in range(len(l))]

```

```

        return pd.DataFrame(yp)

    def score(self, y_true, y_predict):
        return accuracy_score(y_true, y_predict)

    def confusion_matrix(self, y_true, y_predict):
        return confusion_matrix(y_true, y_predict)

```

#### 4. Verify the model

In [ ]: `from sklearn.metrics import ConfusionMatrixDisplay`

```

def visualize_test(clf:AgeGroupPredictorKNN,cm):
    fig, ax =plt.subplots(1,2,figsize = (16, 6))
    corr = pd.DataFrame(clf.X_pca).corr()
    sns.heatmap(corr, ax=ax[0])
    ax[0].set_title("Correlation Matrix")
    ax[1].set_title("Confusion matrix")
    disp = ConfusionMatrixDisplay(cm)
    disp.plot(ax=ax[1])
    plt.show()

```

In [ ]: `import seaborn as sns`

```

params = [{
    'n_components':5,
    'n_init':10,
    'max_iter':1000,
}, {
    'n_components':5,
    'n_init':100,
    'max_iter':2000,
}, {
    'n_components':3,
    'n_init':500,
    'max_iter':5000,
}, {
    'n_components':6,
    'n_init':100,
    'max_iter':3000,
}, {
    'n_components':4,
    'n_init':100,
    'max_iter':3000,
},
]

def test(dt):
    for p in params:
        clf = AgeGroupPredictorKNN().with_params(p)
        clf.fit(dt['train'],dt['ytr'])
        pred = clf.predict(dt['test'])
        score = clf.score(dt['yte'],pred)

```



```

cm = clf.confusion_matrix(dt['yte'],pred)
print("=====
print("Parameters:",p, "Score:",score, "\nConfusion matrix:\n",cm)
visualize_test(clf,cm)
print("=====

d = [{
    'train':X_train_m,
    'ytr':y_train_m_group,
    'test':X_test_m,
    'yte':y_test_m_group
},{
    'train':X_train_f,
    'ytr':y_train_f_group,
    'test':X_test_f,
    'yte':y_test_f_group
},
]
print("Prediction for males:")
test(d[0])
print("Prediction for females:")
test(d[1])

```

Prediction for males:



=====  
Parameters: {'n\_components': 5, 'n\_init': 10, 'max\_iter': 1000} Score: 0.33769  
6335078534

Confusion matrix:

```

[[36 31  7  0]
 [23 51 60  7]
 [13 31 31 13]
 [ 6 35 27 11]]



```



Parameters: {'n\_components': 5, 'n\_init': 100, 'max\_iter': 2000} Score: 0.3376  
96335078534

Confusion matrix:



```
[[36 31  7  0]
 [23 51 60  7]
 [13 31 31 13]
 [ 6 36 26 11]]
```



Parameters: {'n\_components': 3, 'n\_init': 500, 'max\_iter': 5000} Score: 0.3376  
96335078534

Confusion matrix:



```
[[36 31  7  0]
 [23 51 60  7]
 [13 31 31 13]
 [ 6 36 26 11]]
```



Parameters: {'n\_components': 6, 'n\_init': 100, 'max\_iter': 3000} Score: 0.3376  
96335078534

Confusion matrix:



```
[[36 31  7  0]
 [23 51 60  7]
 [13 31 31 13]
 [ 6 36 26 11]]
```



Parameters: {'n\_components': 4, 'n\_init': 100, 'max\_iter': 3000} Score: 0.3376  
96335078534

Confusion matrix:

```
[[36 31  7  0]
 [23 51 60  7]
 [13 31 31 13]
 [ 6 35 27 11]]
```





=====  
Prediction for females:  
=====

Parameters: {'n\_components': 5, 'n\_init': 10, 'max\_iter': 1000} Score: 0.33027  
52293577982

Confusion matrix:


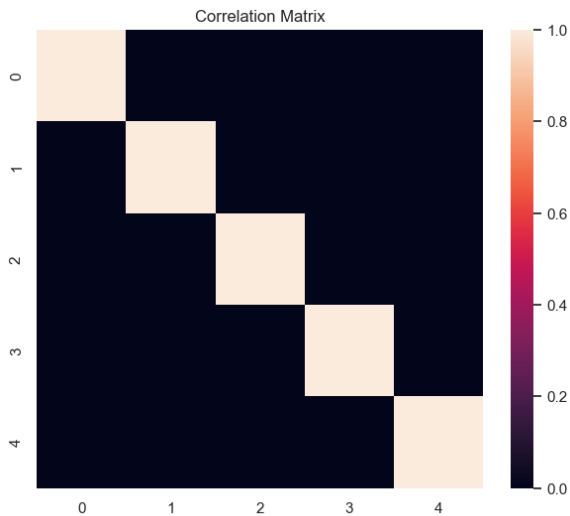
```
[[16  9  0 19]
 [26 44 11 35]
 [ 8 26 22 27]
 [10 31 17 26]]
```



=====  
Parameters: {'n\_components': 5, 'n\_init': 100, 'max\_iter': 2000} Score: 0.3302  
752293577982

Confusion matrix:



```
[[16  9  0 19]
 [26 44 11 35]
 [ 8 26 22 27]
 [10 31 17 26]]
```



Parameters: {'n\_components': 3, 'n\_init': 500, 'max\_iter': 5000} Score: 0.3302  
752293577982

Confusion matrix:



```
[[16  9  0 19]
 [26 44 11 35]
 [ 8 26 22 27]
 [10 31 17 26]]
```



Parameters: {'n\_components': 6, 'n\_init': 100, 'max\_iter': 3000} Score: 0.3302  
752293577982

Confusion matrix:



```
[[16  9  0 19]
 [26 44 11 35]
 [ 8 26 22 27]
 [10 31 17 26]]
```



```
Parameters: {'n_components': 4, 'n_init': 100, 'max_iter': 3000} Score: 0.3302752293577982
```

Confusion matrix:

```
[[16  9  0 19]
 [26 44 11 35]
 [ 8 26 22 27]
 [10 31 17 26]]
```



From the testing result, the best score is about 32% and 33% for male and female respectively. It seems that KNN is not a good choice. Try with another approach to directly comparing the similarity against the matrix after PCAed.

```
In [ ]: from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.metrics.pairwise import cosine_similarity
class AgeGroupPredictorSim:
    def __init__(self):
        self.n_components = 5
        self.max_iter = 10000
        self.n_init = 100

        self.pca = None
        self.X_pca = None
        self.label2group = []
    def with_params(self, params):
        if params == None:
            return self
        for k in params:
            match k:
                case 'n_components':
                    self.n_components = params[k]
                case 'n_init':
                    self.n_init = params[k]
                case 'max_iter':
                    self.max_iter = params[k]
```

```

        return self

    def fit(self,X,y):
        self.pca = PCA(n_components=self.n_components).fit(X)
        self.X_pca = self.pca.transform(X)

        n_clusters = len(y.unique())
        self.knn = KMeans(
            n_clusters=n_clusters,
            n_init=self.n_init,
            max_iter=self.max_iter,
        ).fit(self.X_pca)

        self.map_label_to_group(y)

    def map_label_to_group(self,y):
        self.label2group = y

    def predict(self, X):
        x_pca = self.pca.transform(X)
        sims = cosine_similarity(self.X_pca,x_pca)
        predicts = [self.label2group.iloc[i] for i in np.argmax(sims, axis = 0)]
        return pd.DataFrame(predicts)

    def score(self, y_true, y_predict):
        return accuracy_score(y_true, y_predict)

    def confusion_matrix(self, y_true, y_predict):
        return confusion_matrix(y_true, y_predict)

```

```

In [ ]: import seaborn as sns
        params = [{
            'n_components':5,
            'n_init':10,
            'max_iter':1000,
        },{
            'n_components':5,
            'n_init':100,
            'max_iter':2000,
        },{
            'n_components':3,
            'n_init':500,
            'max_iter':5000,
        },{
            'n_components':6,
            'n_init':100,
            'max_iter':3000,
        },{
            'n_components':4,
            'n_init':100,
            'max_iter':3000,
        },
        ]

    def test(dt):
        for p in params:
            clf = AgeGroupPredictorSim().with_params(p)

```

```

clf.fit(dt['train'],dt['ytr'])
pred = clf.predict(dt['test'])
score = clf.score(dt['yte'],pred)
cm = clf.confusion_matrix(dt['yte'],pred)
print("=====")
print("Parameters:",p, "Score:",score, "\nConfusion matrix:\n",cm)
visualize_test(clf,cm)
print("=====")

```

```

d = [{
    'train':X_train_m,
    'ytr':y_train_m_group,
    'test':X_test_m,
    'yte':y_test_m_group
},{
    'train':X_train_f,
    'ytr':y_train_f_group,
    'test':X_test_f,
    'yte':y_test_f_group
},
]
print("Prediction for males:")
test(d[0])
print("Prediction for females:")
test(d[1])

```

Prediction for males:



=====  
Parameters: {'n\_components': 5, 'n\_init': 10, 'max\_iter': 1000} Score: 0.4293193717277487

Confusion matrix:

```

[[40 21 10  3]
 [30 66 31 14]
 [ 3 32 26 27]
 [ 4 19 24 32]]

```







Parameters: {'n\_components': 5, 'n\_init': 100, 'max\_iter': 2000} Score: 0.4293  
193717277487

Confusion matrix:



```
[[40 21 10  3]
 [30 66 31 14]
 [ 3 32 26 27]
 [ 4 19 24 32]]
```



Parameters: {'n\_components': 3, 'n\_init': 500, 'max\_iter': 5000} Score: 0.4581  
151832460733

Confusion matrix:



```
[[36 26  8  4]
 [19 71 35 16]
 [10 31 28 19]
 [ 3 22 14 40]]
```



Parameters: {'n\_components': 6, 'n\_init': 100, 'max\_iter': 3000} Score: 0.4188  
48167539267

Confusion matrix:



```
[[33 25 13  3]
 [29 63 33 16]
 [ 5 37 28 18]
 [ 3 20 20 36]]
```



Parameters: {'n\_components': 4, 'n\_init': 100, 'max\_iter': 3000} Score: 0.4188  
48167539267

Confusion matrix:

```
[[37 21 14  2]
 [30 64 31 16]
 [ 7 29 30 22]
 [ 8 17 25 29]]
```





Prediction for females:

Parameters: {'n\_components': 5, 'n\_init': 10, 'max\_iter': 1000} Score: 0.40672782874617736

Confusion matrix:



```
[[16 24  1  3]
 [24 53 24 15]
 [ 4 31 30 18]
 [ 3 27 20 34]]
```



Parameters: {'n\_components': 5, 'n\_init': 100, 'max\_iter': 2000} Score: 0.40672782874617736

Confusion matrix:



```
[[16 24  1  3]
 [24 53 24 15]
 [ 4 31 30 18]
 [ 3 27 20 34]]
```



Parameters: {'n\_components': 3, 'n\_init': 500, 'max\_iter': 5000} Score: 0.3211  
009174311927

Confusion matrix:



```
[[11 16 15  2]
 [20 47 30 19]
 [ 9 36 20 18]
 [ 9 25 23 27]]
```



Parameters: {'n\_components': 6, 'n\_init': 100, 'max\_iter': 3000} Score: 0.3822  
62996941896

Confusion matrix:



```
[[17 21  4  2]
 [19 47 32 18]
 [ 7 29 28 19]
 [ 3 25 23 33]]
```



```
Parameters: {'n_components': 4, 'n_init': 100, 'max_iter': 3000} Score: 0.3883
7920489296635
```

Confusion matrix:

```
[[13 25  3  3]
 [24 53 27 12]
 [ 6 35 23 19]
 [ 3 27 16 38]]
```



With the new model which is similarity+PCA, the result improved. Now the best score for male increased to 46%, while female increased slightly to 35%.

## Conclusion

Although the best score is increased to 46% with similarity based solution, the overall result with different models and also several parameters is still not good enough. Here are the possible reasons for the bad result:

- Data is insufficient.
- Biometrics is also not sufficient, esp. considering the correlation between the features are high.
- It is also possible that there is not any relationship between the measured biometrics and age.

To further verify whether it is possible to predict, more data is needed, and also it needs to be thought about to add more features.